1 1. Introduction

How does Federal Reserve communication move asset prices, and when does it move the portfolios of households and institutions? The question is old; our answer is an identification. We decompose monetary-policy news into a conventional component (MP) and an information component (CBI) by letting financial markets invert the central bank's statement — the same-day market response, read through the Jarociński-Karadi sign structure, *is* the decomposition. The sign restriction is not an assumption layered on top of the data; it is the inversion rule that rational expectations implies, stated as a one-line classification on the shock file.

The theoretical anchor. Under rational expectations, the announcement-day price response is the market's Bayesian posterior update given the statement. It follows that any analyst-side measure of what the statement "said" — dictionary scores, LLM scores, human readings — measures the *text*, not the *market's information*. Pillar 1 disciplines everything that follows: text measures and inversion measures must be kept separate, and a text measure that correlates with the market's inversion outcome is suspect by construction. Pillar 2 supplies the decomposition tradition (Kuttner's window; Bauer-Swanson's oracle; Jarociński-Karadi's sign identification) and locates our contribution: statements, decomposed, with the one-line rule. Pillar 3 adds the portfolio side — a risk-ladder model in which information shocks reprice risk premia and flows reallocate along a 1–7 asset-class risk rank.

What we find. 1. *Prices* react to both components, and the dictionary-measured *tone* of the statement matters for prices only in the forward-guidance era — on equities: CRSP VW $\beta = -1.119$ ($p = 0.004$), S&P 500 $\beta = -1.027$ ($p = 0.006$), tone-only $R^2 = 14\%$; in a horse race with both shocks, all three regressors are significant and the R^2 reaches 48%; on 10-year futures the FG interaction is positive (+0.050, $p = 0.030$) — hawks lift yields and sink valuations simultaneously. Outside FG, tone is dead on every asset. 2. *Portfolios* do not: neither shocks nor tone nor LLM-scored text add explanatory power to monthly fund flows (H4: $p = 0.64$ raw; $p = 0.17$ LLM). This null is the paper's ruling result — "prices first, portfolios later." 3. The risk-taking channel runs on *hawkish information*: positive CBI shocks — the market reading the announcement as the central bank acting on information it holds, typically bad news about the economy — drive de-risking along the risk ladder (Risk × CBI = −0.066, $p = 0.036$; sign-asymmetric: Risk × CBI(+) = −0.088, $p = 0.020$ vs Risk × CBI(−) = −0.055, $p = 0.10$; government-bond CBI flow beta = +0.76, $p = 0.037$). Portfolios respond to what the Fed *knows*, not to what it *does* — with a zero-rate-regime exception (MP × ZLB = −1.505, $p = 0.005$) that reveals the conventional channel working through rescue capacity. 4. *Measurement*: the LM dictionary fails by compression (not by mis-reading) exactly when the Fed's communication paradigm shifts — on the 11 divergence statements the dictionary output is nearly constant (std 1.10) while multi-model LLM scoring separates directions sharply (std ≈ 70); sign agreement across four vendor families is 100% on those statements. The LLM measure is positioned as a diagnostic of dictionary failure, not a better scorer.

Contribution. Three claims. (i) *Identification*: the inversion rule, stated as a reproducible one-line classification, gives statement-level MP/CBI decomposition without proprietary infrastructure. (ii) *Mechanism*: the flow null is not a limitation but evidence — with a theoretical reason (timing structure: prices are instantaneous, flows are monthly, and what survives to the monthly horizon is risk content, not words) and a reversal pattern in the window decomposition that pins the reallocation dynamics. (iii) *Measurement*: the resolution-loss signature, its five-layer audit, and the cross-vendor confirmation discipline the tone literature

exactly where its objects are moving — the paradigm-shift episodes.

Roadmap. Section 2 develops the three pillars and derives T1–T5. Section 3 documents data and measurement. Sections 4–6 trace the mechanism through prices, flows, and the risk ladder. Section 7 delivers the measurement-failure evidence. Section 8 collects robustness. Section 9 concludes. # 8. Robustness

1.1 8.1 Orthogonalized tone

The expert-facing concern — tone may proxy the shocks — is addressed by orthogonalizing tone to Target+Path before entering it in every specification. The orthogonalization is a near no-op (first-stage R^2 = 0.4%; raw vs orth tone coefficients in the flow regressions: +0.0231 vs +0.0239), and the price-side horse race (Section 4.3) shows tone surviving *jointly* with both shocks, so the FG-regime text effect is not a shock in disguise.

1.2 8.2 Flow windows: the null and its dynamics

Table 8.1 splits the flow response by window. Same-month government-bond flows carry the de-risking inflow (CBI +0.759, p = 0.037); the *difference* window (post minus same) reverses it (CBI −0.988, p = 0.014). MP is a precise zero in all three windows ($p \geq 0.76$). Two readings, both consistent with the mechanism: households receive the information shock into safe assets first, then rebalance — a reversal pattern, not a drift; and the conventional component never moves flows at any horizon.

1.3 8.3 Cross-measure nulls

The flow null is not tied to one text measure: dictionary tone (p = 0.64), its orthogonalized version (p = 0.63), the FG-regime spec (p = 0.94), and the audited LLM score (p = 0.17) all fail to enter. The price-side null outside FG likewise holds for every asset in Panel B of Table 4.1.

1.4 8.4 Pipeline audit

Every headline number runs through the archival pipeline: the shock file (single CSV, versioned), the dataset builder with sanity checks (statement length, tone–length correlation, CRSP same-day cross-checks), and per-chapter result JSONs committed to the repository. Statement scores carry a five-layer audit (Appendix A) with per-statement verdicts; the two drafting-stage corrections (Operation Twist's router one-off; retirement of two untraceable recalled numbers) are documented in the audit notes rather than silently fixed. # 9. Conclusion

We asked how Federal Reserve communication reaches asset prices and household portfolios, and answered with an identification — the market's inversion rule — rather than a new text score. The answer has three parts. On the announcement day, both the conventional and the information components of policy news move prices with the Jarociński-Karadi sign structure, and the statement's tone carries incremental price information precisely in the forward-guidance era — hawks lift yields and sink valuations at once, and the text survives a horse race against both shocks. Monthly flows ignore the words entirely: the null holds for the shocks, for tone, and for an audited LLM measure, at every window we can construct. What flows do respond

to is the *risk content* of the information — a de-risking ladder on hawkish information, sign-asymmetric, with a zero-rate-regime exception that reveals the conventional channel operating through rescue capacity.

For the communication-paradigm literature, the measurement chapter delivers a caution with teeth: the dictionary does not mis-read paradigm-shift statements, it fails to differentiate them (65× resolution loss on the eleven episodes that define the shifts), so dictionary-based tone findings on exactly those episodes are under-identified. The LLM-based diagnostic — audited, cross-vendor confirmed, and positioned as a measurement tool rather than a competitor — is the repair we offer.

Limitations. Flow aggregation is monthly and class-level; the window structure, however carefully bounded, cannot separate intra-month timing. The text measures cover statements only — minutes and speeches are outside scope. The LLM layer inherits vendor dependence; the committee rule and the audit protocol bound, but do not eliminate, that dependence.

Policy reading. If prices clear the words and portfolios follow the risk, then central-bank communication works on household portfolios *through repricing*, not through persuasion — and communication regimes that shift the paradigm (forward guidance, structural operations) are precisely where text-measurement tools must be re-validated before their scores are interpreted.